

Legged Robots Practical: Project 2

14.11.2023

Plan

Single
Leg

- **W1:** Introduction
- **W2:** Derivation of double pendulum's kinematics and dynamics
- **W3:** Jacobian (Cartesian PD + Force Control)
- **W4:** Inverse Kinematics (compare with force control)
- **W5:** Single-leg hopping

Biped


- **W6-8:** Biped Walking with Divergent Component of Motion (BD Atlas)
[Mini-Report]

Quadruped

- **W9-10:** Quadruped CPG Trot
- **W11-14:** Quadruped Locomotion Project (CPGs, Deep RL)
[Report - Quadruped] [COMPETITION 19.12.2023]

Part 1: Central Pattern Generators

Trot




Bound






Pace






Walk




Part 2: Deep Reinforcement Learning






Quadruped Model Reference Frame




Quadruped Model Leg References



Quadruped Model Joint References



Joint angles \leftrightarrow Cartesian space (in leg frame)



$$p = f(q)$$

Forward kinematics

$$q = f^{-1}(p)$$

Inverse kinematics


$$\dot{p} = v = J(q)\dot{q}$$

Foot linear velocity

$$\tau = J^T(q)F$$

Map desired end effector force to torques

Joint angles \leftrightarrow Cartesian space (leg frame control)



$$p = f(q)$$

Forward kinematics

$$q = f^{-1}(p)$$

Inverse kinematics

$$\dot{p} = v = J(q)\dot{q}$$

Foot linear velocity

$$\tau = J^T(q)F$$

Map desired end effector force to torques

$$\tau_{joint} = K_{p,joint}(q_d - q) + K_{d,joint}(\dot{q}_d - \dot{q})$$

Joint PD

$$\tau_{Cartesian} = J^T(q)[K_{p,Cartesian}(p_d - p) + K_{d,Cartesian}(v_d - v)]$$

Cartesian PD

$$\tau_{final} = \tau_{joint} + \tau_{Cartesian}$$

Contributions from both joint PD and Cartesian PD

Central Pattern Generators: Review

100%


From Lecture 4

Descending modulation

Spinal cord

Reflexes

Central pattern generators



Musculoskeletal system

Respective Role
in motor control



lamprey



salamander



"Comp ima sp



cat



human

Modeling the CPG with coupled oscillators (Quadruped)

Amplitude:


$$\dot{r}_i = \alpha(\mu - r_i^2)r_i$$

Phase:


$$\dot{\theta}_i = \omega_i + \sum_{j=0}^3 r_j w_{ij} \sin(\theta_j - \theta_i - \phi_{ij})$$

Output:

$$x_{\text{foot}} = -d_{\text{step}} r_i \cos(\theta_i)$$
$$z_{\text{foot}} = \begin{cases} -h + g_c \sin(\theta_i) & \text{if } \sin(\theta_i) > 0 \\ -h + g_p \sin(\theta_i) & \text{otherwise} \end{cases}$$



Mapping CPG States to Foot Positions with Inverse Kinematics



$$\dot{r}_i = \alpha(\mu - r_i^2)r_i$$



$$\dot{\theta}_i = \omega_i$$

$$x_{\text{foot}} = -d_{\text{step}}r_i \cos(\theta_i)$$

$$z_{\text{foot}} = \begin{cases} -h + g_c \sin(\theta_i) & \text{if } \sin(\theta_i) > 0 \\ -h + g_p \sin(\theta_i) & \text{otherwise} \end{cases}$$


Gait Terminology

- *Stride duration* = the duration of a complete cycle (the period)
- *Swing phase* of a limb (period during which the limb is off the ground)
- *Stance phase* (period during which the limb touches the ground)
- *Duty factor* = Stance duration / Stride duration




Most common quadruped gaits


Classification in terms of the footfall sequences (mainly used in mathematical biology)




Lateral sequence walk



Diagonal sequence walk




Trot




Pace

Symmetric


Asymmetric



Bound




Rotary gallop



Transverse gallop

Most common quadruped gaits

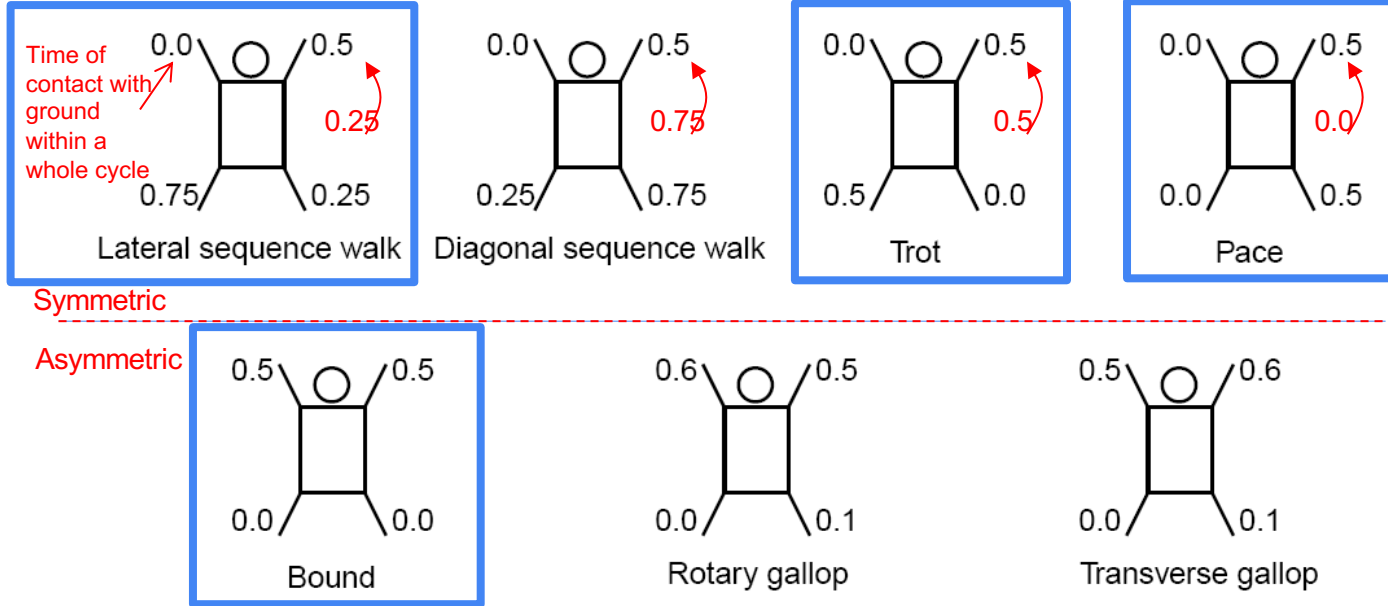
Classification in terms of the footfall sequences (mainly used in mathematical biology)



This project

Most common quadruped gaits

Classification in terms of the footfall sequences (mainly used in mathematical biology)



This project

$$\dot{r}_i = \alpha(\mu - r_i^2)r_i$$

$$\dot{\theta}_i = \omega_i + \sum_{j=0}^3 r_j w_{ij} \sin(\theta_j - \theta_i - \phi_{ij})$$

What should ϕ be for each gait?


Deep Reinforcement Learning: Review

From Lecture 7

Reinforcement Learning

An MDP is defined by:

- Set of states \mathcal{S}
- Set of actions \mathcal{A}
- Transition function $P(s' | s, a)$
- Reward function $R(s, a, s')$
- Start state s_0
- Discount factor γ
- Horizon H



- Return over a trajectory $\tau = (s_0, a_0, s_1, a_1, \dots)$

$$R(\tau) = \sum_{t=0}^{\infty} \gamma^t r_t$$

- Policy $\pi(a_t | s_t)$ maps from states s_t to actions a_t (Goal: find policy maximizing above return)
- Value function: $V^\pi(s) = \mathbb{E}_{\tau \sim \pi}[R(\tau) | s_0 = s]$
- Action-value function: $Q^\pi(s, a) = \mathbb{E}_{\tau \sim \pi}[R(\tau) | s_0 = s, a_0 = a]$
- Advantage function: $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$

Many Existing Tools for Reinforcement Learning

- RL algorithm implementations

- stable-baselines3 <https://github.com/DLR-RM/stable-baselines3>
- ray[rllib] <https://github.com/ray-project/ray>
- spinningup <https://github.com/openai/spinningup>
- tianshou <https://github.com/thu-ml/tianshou/>
- ... many others!

PPO, SAC

- Physics simulators

- pybullet <https://github.com/bulletphysics/bullet3>
- MuJoCo <https://mujoco.org>
- RaiSim <https://raisim.com>
- Isaac-Gym <https://developer.nvidia.com/isaac-gym>
- ... and others!

RL Considerations

Algorithm

- On/off policy
- Hyperparameters
- Network architecture
- Random seeds/trials

...implementation
dependent!

MDP Design Decisions

- Observation space
- Action space
- Reward function

Environment Parameters


- Simulator dynamics
- Control gains –
joint/Cartesian
- Control/environment
time step
- Noise, latency

From Lecture 7

State/Action/Reward Space: A1


s_t ? i.e.

- body (z, r, p, y)
- body velocities
- joint states




r_t ? i.e.

- body linear velocity
- energy penalty



a_t ?

- motor positions/torques
- Cartesian PD
- CPG state modulations





From Lecture 7

State/Action/Reward Space: A1

s_t ? i.e.

- body (z, r, p, y)
- body velocities
- joint states




a_t ?

- motor positions/torques
- Cartesian PD
- CPG state modulations

r_t ? i.e.

- body linear velocity
- energy penalty




This project: construct the MDP


From Lecture 7

Joint Position Control vs. Cartesian PD Control (PPO/SAC)

Action Space: $a_t = q_{1...N}$

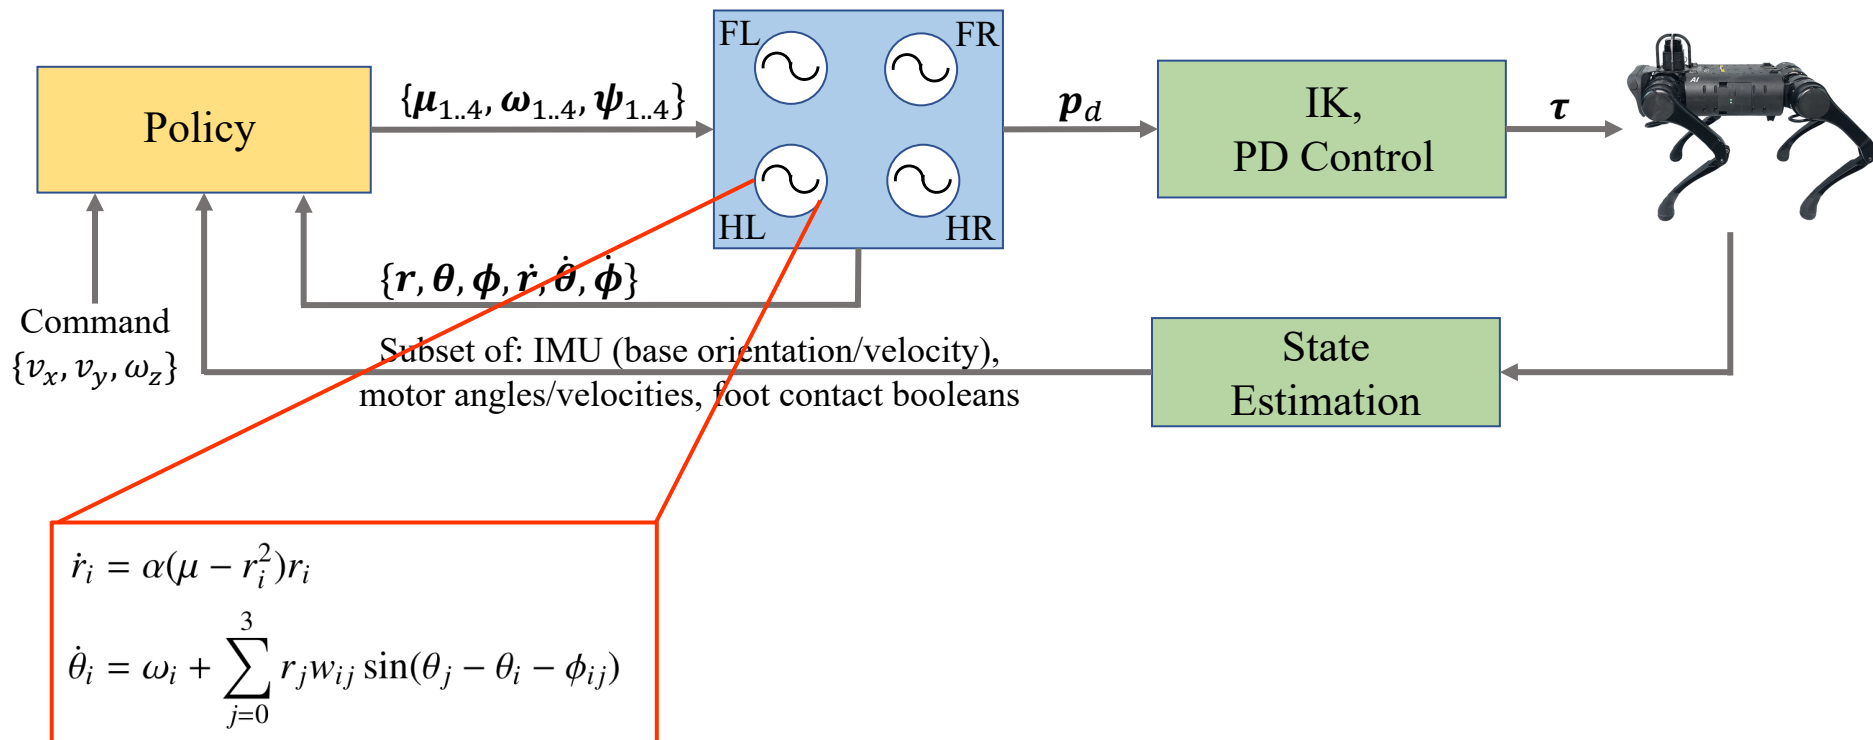


Action Space: $a_t = [x_{ee_i}, y_{ee_i}, z_{ee_i}]$




CPG-RL: Learning Central Pattern Generators for Quadruped Locomotion

From Lecture 7



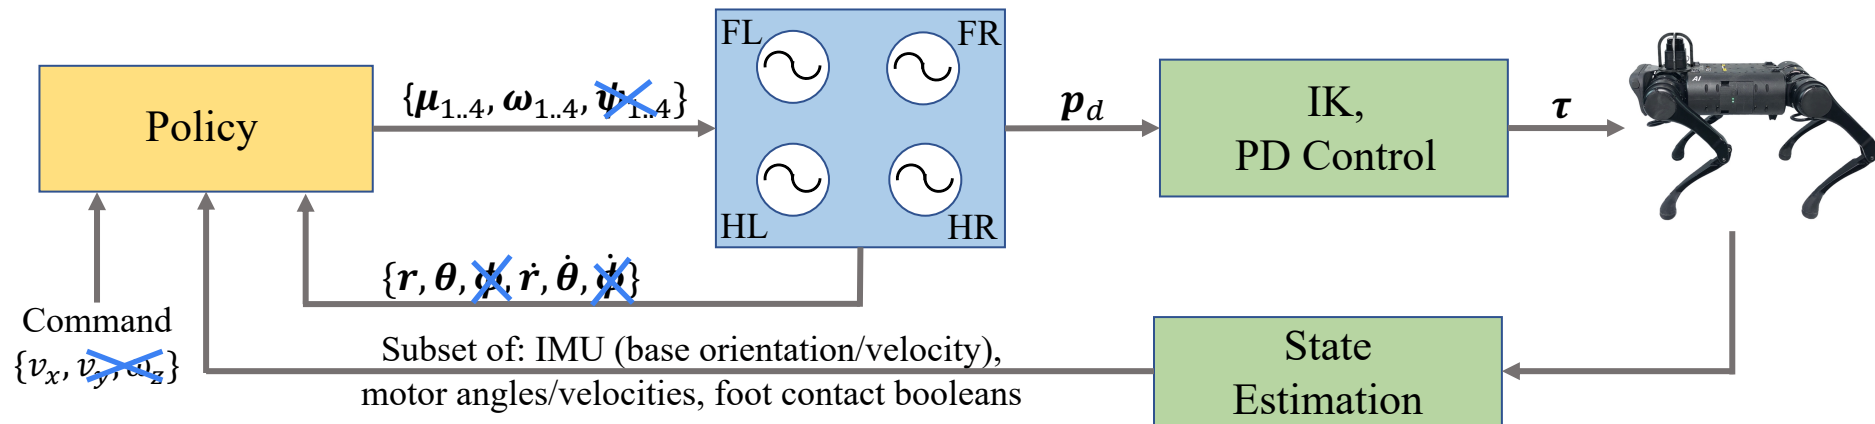
CPG-RL: Learning Central Pattern Generators for Quadruped Locomotion

From Lecture 7




CPG-RL: Learning Central Pattern Generators for Quadruped Locomotion

From Lecture 7





$$\dot{r}_i = \alpha(\mu - r_i^2)r_i$$

$$\dot{\theta}_i = \omega_i + \sum_{j=0}^3 r_j w_{ij} \sin(\theta_j - \theta_i - \phi_{ij})$$



How robust is your approach? To be determined at the 19.12.2023 competition



Goal oriented locomotion: what should be in the observation space and reward function?

Use function:


`get_distance_and_angle_to_goal()`

Sensory information in `quadruped.py`:

- `GetBaseOrientation()`

- `GetBaseLinearVelocity()`

...



Tips

- Monitor episode length and reward mean during training
- Training should complete within 1 million timesteps for reasonable observation space, action space, and reward function choices (with no noise in the environment)
- No training on test environment (used for competition)
- Start training early!